

Greedy Algorithms: Off-line caching (Textbook Problem 16-5)

Modern computers use a cache to store a small amount of data in a fast memory. When a computer program executes, it makes a sequence $\langle r_1, r_2, \dots, r_n \rangle$ of n memory requests, where each request is for a particular data element. For example, a program that accesses 4 distinct elements $\{a, b, c, d\}$ might make the sequence of requests $\langle e, b, d, b, d, a, c, e, b, a, c, b \rangle$. Let k be the size of the cache. When the cache contains k elements and the program requests the $(k + 1)$ st element, the system must decide, for this and each subsequent request, which k elements to keep in the cache. In other words, **cache hit** and **cache miss** will happen. The cache replacement algorithm evicts data with the goal of *minimizing the number of cache misses* over the entire sequence of requests.

Typically, caching is an on-line problem. That is, we have to make decisions about which data to keep in the cache without knowing the future requests. Here, however, we consider the off-line version of this problem, in which we are given in advance the entire sequence of n requests and the cache size k .

We can solve this off-line problem by a greedy strategy called **furthest-in-future**, which chooses to evict the item in the cache whose next access in the request sequence comes furthest in the future.

- (a) Write pseudocode for a cache manager that uses furthest-in-future strategy. Assume that request sequence is of size n and cache size is k .
- (b) What is the running time of this algorithm in terms of n, k .
- (c) Justify its optimality.
Hint: Greedy choice property and optimal substructure property